

Marginal propensities to consume with behavioural agents*

Andrej Mijakovic[†]

European University Institute

October 23, 2022

Abstract

I use survey questions on spending in hypothetical scenarios to study differences in marginal propensities to consume out of positive and negative income shocks. I find a sizeable asymmetry in spending responses driven by large adjustments following income losses compared to small adjustments following income gains. This asymmetry occurs irrespective of the position in the wealth distribution and is not driven by liquidity constraints. As such, it stands in stark contrast to predictions of traditional models of consumption behaviour. I develop a model with dissaving-averse households, a behavioural feature consistent with mental accounting, that rationalizes this asymmetry and simultaneously generates low MPCs out of income news, low MPCs out of wealth and high MPCs out of income gains for households that are not liquidity-constrained. I show through the lens of a quantitative model that asymmetric MPCs dampen the effectiveness of redistributive fiscal policy.

*I am indebted to my advisors Edouard Challe and Russell Cooper for their invaluable guidance and support. I would also like to thank Simon Gilchrist, Tullio Jappelli (discussant) and seminar and conference participants at the European University Institute, New York University, the 4th Behavioral Macroeconomics Workshop, the Annual Conference of the Money, Macro and Finance Society 2022 and the 2022 CEPR European Conference on Household Finance for helpful comments and suggestions.

[†]andrej.mijakovic@eui.eu.

1 Introduction

The marginal propensity to consume (MPC) plays a central role in the transmission of fiscal and monetary policy. In quantifying MPCs, the empirical literature has identified several puzzles that are difficult to reconcile with traditional theories of household consumption behaviour. Empirically observed consumption responses tend to be excessively sensitive to contemporaneous income changes, insensitive to changes in wealth and news about future income, and only loosely related to liquidity constraints.

This paper first aggravates the divide between theory and data by documenting an additional puzzle: consumption responds more strongly to income losses than to income gains irrespective of the degree of liquidity constraints. It then resolves this divide by proposing a unifying theoretical framework of consumption behaviour that addresses several empirical puzzles at once.

I begin by documenting large differences in hypothetical spending responses to positive and negative income shocks for a sample of households from the Fed Survey of Consumer Expectations (SCE). MPCs are asymmetric driven by large adjustments following income losses compared to small adjustments following income gains. The average MPC out of positive income shocks is rather small at 0.2, while the average MPC out of negative income shocks is substantially larger at 0.73. These asymmetries are also widespread. Around 85 percent of households indicate that they would adjust consumption more strongly in response to negative than to positive income shocks.

Notably, consumption responds asymmetrically across all levels of wealth. While the level of the MPC asymmetry, defined as the difference between the MPC out of income losses and income gains, decreases from 0.62 for the bottom quintile to 0.42 for the top quintile of the wealth distribution, this decrease is small relative to the absolute level of the asymmetry. Larger holdings of wealth primarily reduce the MPC out of losses, but insufficiently so to close the gap with the MPC out of gains. These patterns cannot be explained by the presence of wealthy hand-to-mouth households, who hold substantial amounts of wealth but behave like liquidity-constrained households due to a large share of illiquid assets ([Kaplan et al., 2014](#)). Even the most liquid households in the sample display sizeable asymmetries.

Several features of the data corroborate the robustness of the survey evidence on MPCs. First, spending out of the hypothetical scenario is comparable to reported spending out of tax refunds. This direct comparison is possible as survey participants are asked to indicate both their spending response to a hypothetical income change and their actual spending response to a realized tax refund. Second, households' ex-ante planned expenditure predicts their ex-post realized expenditure, suggesting that households form accurate spending plans. Finally, asym-

metric consumption responses are not driven by the presence of financially illiterate households.

While small sign asymmetries in MPCs, particularly for liquidity-constrained households, are consistent with standard models of consumption behaviour (Kaplan and Violante, 2010), large asymmetries for unconstrained households are at odds with the predictions of conventional consumption models. Several extensions have been suggested that can generate large MPCs out of income gains for unconstrained agents (Kueng, 2018; Ilut and Valchev, 2020; Lian, 2020; Laibson et al., 2021; Boutros, 2022), but they fail to generate large asymmetries, i.e. MPCs out of losses that are considerably larger.

In order to rationalize these empirical patterns, I propose an extension of the standard consumption model that incorporates mental accounting preferences (Shefrin and Thaler, 1988; McDowall, 2019). Households hold different mental accounts with regards to their current income and asset position. Funds pertaining to the asset account are not perfect substitutes for funds in the income account. In particular, consuming out of the asset account is costly, with the cost governed by a parameter that indicates the level of dissaving-aversion.

The partition between the mental accounts for income and assets is given by a savings rule against which households benchmark their savings decisions. This partition delivers a flexible distinction between the income and the asset account and allows, for example, regular retirement savings to fall under the asset account. In the model, I assume that savings rules are formed endogenously by households and adjust flexibly to changes in wealth, but only imperfectly to changes in income. The combination of dissaving-aversion and rigidity in savings rules generates rich non-linearities in the consumption response to income shocks. Empirically, the presence of savings rules is strongly backed by the data. A large majority of households in the SCE sample states to plan their savings.

I show analytically that the mental accounts model generates predictions that are in line with several empirically documented patterns. First, it generates MPCs out of income losses that are larger than MPCs out of income gains. Second, it predicts responses to news about income changes that are smaller than responses to contemporary income changes in excess of what the standard model would predict. Third, it predicts MPCs out of wealth that are lower than MPCs out of income.

Heterogeneity in the presence of savings plans across households in the data furthermore allows me to test the prediction that households with a savings plan display higher MPC asymmetries than households without a savings plan. I provide suggestive evidence that this is supported by the data. Households with a savings plan have MPCs that are significantly more asymmetric.

To study the implications of dissaving-aversion in a quantitative setup, I incorporate mental accounting preferences into a life-cycle model with income risk and borrowing constraints. The

model generates large MPC asymmetries that closely match the ones observed in the data. In particular, it predicts large MPCs out of losses across all levels of wealth while keeping MPCs out of gains moderate. Importantly, neither the asymmetry nor the relation between MPCs and wealth are explicitly targeted in the calibration exercise. The dissaving-aversion motive is only disciplined by matching the average MPC out of losses and shares of households with a savings plan in the SCE.

A comparison with a model without mental accounting preferences shows that the mental accounting model performs well across several dimensions. First, the standard model cannot generate asymmetric MPCs for households that are not liquidity-constrained. Second, the mental accounting model predicts relatively large MPCs out of gains for unconstrained households (Lewis et al., 2019; Fagereng et al., 2021). Third, and in line with the predictions of the analytical model, the mental accounting model generates lower MPCs out of income news and wealth (McDowall, 2019; Di Maggio et al., 2020; Chodorow-Reich et al., 2021; Christelis et al., 2021; Fuster et al., 2021). Crucially, this does not come at the cost of missing other moments of the data. The mental accounting model preserves the dynamics of the frictionless model with respect to the life-cycle profiles and dispersion of consumption and savings.

The policy implications of widespread MPC asymmetries are considerable. They suggest a cautious approach to redistributive fiscal policies as the traditional logic of redistributing from the rich (low MPC) to the poor (high MPC) does not necessarily hold in this framework. Because MPCs out of losses are large even for the rich, the increase in consumption by the poor might not be enough to offset the reduction in consumption by the rich. A simulation of the mental accounting model suggests that the effectiveness of a simple redistributive policy in which the bottom half of the income distribution receives transfers crucially depends on how the transfers are financed. A one-off income tax on the richest quarter of households slightly *reduces* aggregate consumption while taxing wealth instead of income substantially increases aggregate consumption due to low MPCs out of wealth, but high MPCs out of income.

Apart from purely redistributive policies, asymmetric MPCs can also have more general implications. If MPCs at the micro-level are indicative of MPCs at the macro-level, fiscal contractions could translate into stronger consumption responses than fiscal expansions. This is particularly relevant in the light of recent empirical evidence showing asymmetric responses to fiscal policy at the aggregate level (Barnichon et al., 2022).

The remainder of this section relates this paper to the literature. Section II describes the data. Section III presents empirical results. Section IV introduces the analytical framework. Section V presents the quantitative model and the fiscal experiment. Section VI concludes.

Literature. This paper is closely related to the empirical literature on MPC asymmetries. Christelis et al. (2019) and Fuster et al. (2021) document that consumption responds more

strongly to negative than to positive income shocks using survey responses to hypothetical income changes. [Bracha and Cooper \(2014\)](#), [Sahm et al. \(2015\)](#) and [Bunn et al. \(2018\)](#) provide similar evidence from reported consumption responses to actual income changes. [Baugh et al. \(2021\)](#) finds asymmetric responses to expected income changes using transaction-level data. This paper contributes to the literature by studying the effects of a relatively large hypothetical income shock and showing that asymmetric consumption responses are sizeable across all levels of wealth and liquidity. Differently to other studies, this paper also provides a theoretical framework that generates large asymmetries for households which are not liquidity-constrained.

The theoretical framework outlined in this article connects to a wider literature that finds consumption patterns that are inconsistent with predictions of traditional consumption models. First, MPCs out of income gains can be large even for households that are not liquidity-constrained ([Lewis et al., 2019](#); [Fagereng et al., 2021](#)). Second, several studies suggest that MPCs out of wealth are lower than MPCs out of income ([Di Maggio et al., 2020](#); [Chodorow-Reich et al., 2021](#); [Christelis et al., 2021](#)). Third, there is ample evidence that consumption is insensitive to the receipt of news about future income, but excessively sensitive once the predictable income change materializes ([Kueng, 2018](#); [Olafsson and Pagel, 2018](#); [Ganong and Noel, 2019](#); [McDowall, 2019](#); [Fuster et al., 2021](#)).

Several studies propose behavioural extensions that rationalize selected aspects of the empirically observed consumption responses. Most of these extensions are focused on generating high MPCs out of income gains. Prominent examples include present bias ([Laibson et al., 2021](#)), temptation preferences ([Attanasio et al., 2020](#)), near-rationality ([Kueng, 2018](#)), bounded rationality ([Boutros, 2022](#)), anticipation-dependence ([Thakral and Tô, 2021](#)) or imperfect reasoning ([Ilut and Valchev, 2020](#)). [Lian \(2020\)](#) proposes a mechanism through which anticipation of future mistakes amplifies both MPCs out of gains and losses. [Ganong and Noel \(2019\)](#) introduce households with present bias to generate high MPCs out of predictable income losses. Most closely related to the model presented here, [McDowall \(2019\)](#) introduces a mental accounting friction to explain high MPCs out of predictable income gains. In contrast to these studies, my model explicitly addresses MPC asymmetries. At the same time, it is also consistent with the previously mentioned consumption patterns.

2 Data

I measure MPCs using hypothetical survey questions from the New York FED Survey of Consumer Expectations (SCE). The SCE is a monthly online survey of a rotating panel of around 1,300 households. It collects information on household expectations and decisions on a broad variety of topics and provides detailed accounts of household income, balance sheets and demographics. As such, it covers a wide range of variables that are typically considered to affect

MPCs.

Survey questions about hypothetical scenarios are widely used to elicit MPCs (Jappelli and Pistaferri, 2014; Bunn et al., 2018; Christelis et al., 2019; Fuster et al., 2021) and offer several advantages compared to other methods. First, they provide a simple way to measure MPCs out of negative income shocks. Other approaches, such as quasi-natural experiments (Parker et al., 2013; Fagereng et al., 2021) or semi-structural methods (Blundell et al., 2008; Kaplan et al., 2014; Commault, 2020), are often limited to the measurement of MPCs out of positive income shocks or a mix of positive and negative shocks. Second, they measure both MPCs out of positive *and* negative income shocks for the same household at the same point in time. This is important if households differ structurally in the types of shocks they face or if MPCs vary over time, for example over the business cycle. Third, the survey format allows me to study the same income shock for all households with respect to its magnitude. Other methods frequently average over various shock sizes, even though the magnitude of the shock both theoretically and empirically affects the level of the MPCs.

One might suspect that households' actual consumption choices differ from their intended consumption choices, as stated for example in response to hypothetical scenarios. The literature suggests that MPCs are quite robust to the choice of measurement. Within the context of the 2008 and 2020 stimulus payments, Parker and Souleles (2019) and Parker et al. (2021) compare self-reported consumption responses with actual consumption responses and find that the reported use of stimulus payments is highly informative about the household's actual spending response. Bunn et al. (2018) compares reported MPCs to MPCs from hypothetical survey questions and finds similar values. Shapiro and Slemrod (2003) and Sahm et al. (2010) find that ex-ante intended and ex-post reported consumption responses are comparable, while Graziani et al. (2016) finds ex-post consumption responses that are larger than originally intended. Sahm et al. (2015) finds that such responses are particularly aligned for tax increases, i.e. negative income shocks.

2.1 MPC measure

The SCE directly measures MPCs through the following two questions:

Suppose next year you were to find your household with 10% more income than you currently expect. What would you do with the extra income?

Now imagine that next year you were to find yourself with 10% less household income. What would you do?

Participants are asked to give both a qualitative and a quantitative response in which they specify what percentage of additional income they would spend, save or use to pay down debt or,

in the case of income loss, what percentage would be absorbed by reducing spending, depleting savings or borrowing.¹

Some caveats apply to the phrasing of the questions and response options, which are ambiguous along some dimensions. The term *spending* could refer to both non-durable and durable consumption. As such, I remain agnostic on which type of consumption the MPC measure is capturing. It could equivalently be interpreted as a marginal propensity to spend (MPX) (Laibson et al., 2022). The question is also vague about the horizon over which households would increase or decrease their spending. Lastly, households might have different interpretations regarding the persistence of the income shock. In order to map the empirical MPC to the theoretical framework, I will assume that households interpret the income change as transitory. This is supported by the fact that the level of the MPC out of income gains is comparable to the level found in other articles studying transitory income changes, as discussed in the next section. In general, as long as the same household interprets the two (identically phrased) questions regarding income gains and losses in the same way, the difference or asymmetry between consumption responses to positive and negative shocks should be captured adequately, even if the level of the individual MPC measures is biased.²

2.2 Sample description

I combine the monthly SCE core survey with two additional modules at lower frequency, the Spending Survey and the Household Finance Survey. Incorporating information on household balance sheets from the Household Finance Survey comes at the expense of losing the panel dimension.³ I only keep households which respond to both MPC questions. Lastly, I winsorize balance sheet variables at the 1 percent level. This yields a cross-section of 4,009 households over the period 2015-2018.

Table 1 reports summary statistics for demographic and financial variables. It also reports summary statistics for MPCs, which are discussed in the next section. The SCE is designed to be nationally representative, but it somewhat oversamples higher income, wealthier and more educated households. It provides survey weights to account for these differences. A detailed comparison of SCE data with the American Community Survey and Survey of Consumer Finances can be found in Fuster et al. (2021).

¹Appendix A provides more detailed information on the qualitative response options.

²Despite the identical phrasing of the questions, one might be concerned that households interpret the persistence of the gain and loss scenarios differently due to differently experienced histories of shocks. Empirical evidence, however, suggests that the persistence of positive and negative income shocks is similar for the median household (Arellano et al., 2017; Guvenen et al., 2021).

³MPCs are similar between the larger sample without balance sheet variables and the final sample. The panel dimension of the larger sample also allows me to establish that MPCs are fairly stable within households.

Table 1: Summary statistics

	Mean	Median	Std. dev.	Min	Max	N
<i>Demographics</i>						
Age	50.72	51.00	15.24	18	96	4,009
Female	0.48	0.00	0.50	0	1	4,009
College degree	0.56	1.00	0.50	0	1	4,009
Homeowner	0.74	1.00	0.44	0	1	3,684
<i>Financial variables</i>						
Income	82,139	65,000	69,547	150	400,000	3,630
Bank holdings	28,348	8,000	69,363	0	1,600,000	2,623
Liquid assets	90,409	10,000	234,445	0	1,600,000	3,450
Liquid debt	27,695	10,000	48,463	0	300,000	3,660
Total assets	450,130	239,000	602,383	0	4,585,000	3,284
Total debt	96,766	36,500	133,111	0	880,000	3,642
<i>Spending responses</i>						
MPC+	0.20	0.10	0.24	0	1	4,009
MPC-	0.73	0.85	0.31	0	1	4,009

Notes: Liquid assets include money in checking and savings accounts, stocks and bonds. Total assets additionally include retirement funds and housing wealth. Liquid debt includes credit card debt, auto loans, student loans, and medical or legal bills. Total debt additionally includes mortgages.

3 Empirical results


This section presents the main results of the empirical analysis. I first show cross-sectional evidence on MPC asymmetries. I then illustrate that MPC asymmetries are only weakly related to observable characteristics and are present irrespective of the household’s position in the wealth distribution. Finally, I present several robustness checks corroborating the validity of the survey data.

3.1 MPC asymmetries

Figure 1 shows the distribution of MPCs across households for both income gains (MPC^+) and income losses (MPC^-). The average MPC^+ is 0.2, which is at the lower end of empirical estimates (see for example the review in [Jappelli and Pistaferri \(2010\)](#)), but in line with the notion that larger income windfalls induce lower relative consumption responses. Almost half of the households indicate little to no consumption adjustment in response to a positive income shock, while only a negligible share indicates to spend all additional income.


This pattern flips completely for MPCs out of income losses. Almost half of the households fully absorb the income loss by cutting consumption, with only few households not adjusting

Figure 1: Distribution of MPCs out of income gains and income losses



Notes: Dashed lines denote the average MPC^+ and MPC^- in the sample.

Figure 2: Distribution of MPC asymmetries



Notes: The dashed line denotes the average MPC asymmetry in the sample.

consumption in response to the loss. The average MPC out of income losses is substantial at 0.73. This value is comparable to the values found in [Bunn et al. \(2018\)](#) and [Bracha and Cooper \(2014\)](#), but larger than in [Christelis et al. \(2019\)](#), [Surico and Trezzi \(2019\)](#) or [Fuster et al. \(2021\)](#).

To emphasize this divergence, Figure 2 plots the distribution of MPC asymmetries, defined as the difference between MPC^- and MPC^+ for each household. Almost all households adjust consumption more strongly to negative than to positive income shocks. Moreover, the asymmetry is quantitatively large. The average asymmetry amounts to 0.53, and a quarter of households fully cuts consumption in response to negative income shocks, but does not increase consumption at all in response to positive income shocks.


3.2 Heterogeneity across the wealth distribution

The existence of sign asymmetries is by itself not surprising. A standard consumption model with borrowing constraints, for example, predicts asymmetric MPCs for borrowing-constrained households. The size and ubiquity of the asymmetry, however, suggest that liquidity constraints cannot be the main driver of this asymmetry.

To understand the role of liquidity, Figure 3 plots the average MPC asymmetry across quintiles of the net liquid wealth distribution.⁴ The asymmetry is present irrespective of the position in the wealth distribution. It decreases in wealth, but only marginally compared to the


⁴Net liquid wealth is defined as the sum of bank deposits, stocks and bonds minus liquid debt, i.e. most types of debt excl. mortgages.

Figure 3: MPCs across the net liquid wealth distribution



Notes: MPC asymmetry is defined as the difference between the MPC out of losses and the MPC out of gains. Grey bars indicate 95% confidence bands.

Figure 4: MPC asymmetries across measures of liquidity constraints



Notes: Each line corresponds to the MPC asymmetry across the respective distribution. Grey bars indicate 95% confidence bands. Total net wealth is defined as total assets (liquid assets + retirement funds and housing wealth) - total debt (liquid debt + mortgages). Bank holdings refer to money in checking and savings accounts.

absolute level of the asymmetry, and much less than theory would predict. While a standard consumption model would predict perfectly symmetric MPCs, the data suggest that even the wealthiest 20 percent have an asymmetry of 0.42.⁵

Figure 3 also decomposes the asymmetry into MPCs out of gains and losses. While MPCs out of gains are rather constant across wealth levels, MPCs out of losses are decreasing in wealth.⁶ As such, the MPC out of losses is the primary driver behind the narrowing of the asymmetry for higher levels of wealth. Still, even the wealthiest households display large consumption responses to negative income shocks.

Irrespective of which wealth measure one looks at, consumption always responds more strongly to income losses than to gains (Figure 4). MPCs are similar across the distributions of liquid and total wealth, i.e. the sum of liquid and illiquid wealth. This speaks against the presence of wealthy hand-to-mouth households driving the results (Kaplan and Violante, 2014).

⁵Figures B1 and B2 show that this asymmetry is also present for the top ten and five percent of the wealth distribution.

⁶The finding that the relation between MPC^+ and wealth is weak at best has been documented before (Bunn et al., 2018; Christelis et al., 2019; Fuster et al., 2021). Low MPCs out of gains for households with little wealth are primarily driven by households that hold net debt. These households predominantly use the income windfall to repay debt, as for example studied in Boutos (2019). Once one excludes net debtors, the relation between MPCs out of gains and wealth is essentially flat, see Figure B3. The positive relation between MPC^+ and wealth could also be due to the large size of the transfer. Andreolli and Surico (2021) detects a similar relation in the data for large, but not small transfers and explains this finding through a model with non-homothetic preferences.

In particular, even when I restrict the wealth definition to only include funds in checking and savings accounts, arguably the most liquid assets apart from cash, MPCs are still highly asymmetric. Finally, Figure 4 shows that also households with substantial liquid wealth relative to income do not smooth consumption in response to income losses. The top quintile holds liquid wealth in excess of annual income and should in theory be able to buffer a loss that amounts to only a small fraction of that.

3.3 Heterogeneity across other dimensions

MPC asymmetries could in principle be associated with observable characteristics other than wealth such as income, homeownership status or age. Figures B4, B5 and B6 show graphically that MPC asymmetries persist nevertheless across these dimensions. To study the relation between MPCs and individual characteristics more formally, I estimate the following specification:

$$MPC_i^j = \beta_0 + \beta_1 wealth_i + \gamma X_i + u_i$$

where MPC_i^j denotes the MPC measure $j \in \{+, -, asymmetry\}$ for household i , $wealth$ is a measure of net liquid wealth and X is a vector of control variables that are typically considered to affect MPCs.⁷

Table 2 shows that observable characteristics only explain a small share of variation in MPCs. As noted earlier, MPC asymmetries are negatively correlated with net liquid wealth (Column 1 and 2). This difference primarily stems from the negative correlation of MPCs out of losses with wealth (Column 5 and 6), but also partly the positive correlation between MPCs out of gains and wealth (Column 3 and 4). Columns 2, 4 and 6 add controls. Older households have somewhat lower MPC asymmetries, as well as households with higher income. Households with mortgages have higher MPC asymmetries due to a lower MPC out of income gains. Income expectations do not seem to significantly affect MPCs.

3.4 Validity

The magnitude and ubiquity of MPC asymmetries might appear surprising to some readers, even though, as previously discussed, hypothetical survey questions have been shown to capture fairly well actual consumption responses. In this section, I conduct three additional exercises to assess the validity of the survey data and corroborate my empirical findings.

MPCs out of hypothetical income gains vs tax refunds: The SCE asks participants how much of their annual tax refund they spent or they planned to spend. This allows me

⁷While the data have a time dimension, I do not control for time effects in the main specification due to negligible variation over time.

Table 2: MPCs and household characteristics

	(1) MPC^{Asy}	(2) MPC^{Asy}	(3) MPC^+	(4) MPC^+	(5) MPC^-	(6) MPC^-
Net liq. asset quintile 2	-0.050** (0.025)	-0.047* (0.026)	0.024* (0.014)	0.012 (0.014)	-0.026 (0.020)	-0.035* (0.020)
Net liq. asset quintile 3	-0.033 (0.025)	-0.031 (0.027)	0.046*** (0.015)	0.025 (0.016)	0.013 (0.019)	-0.005 (0.020)
Net liq. asset quintile 4	-0.160*** (0.026)	-0.145*** (0.027)	0.078*** (0.015)	0.072*** (0.015)	-0.082*** (0.020)	-0.073*** (0.021)
Net liq. asset quintile 5	-0.193*** (0.027)	-0.160*** (0.029)	0.080*** (0.015)	0.066*** (0.016)	-0.113*** (0.022)	-0.094*** (0.023)
Age between 35-55		0.030 (0.023)		0.006 (0.014)		0.035* (0.018)
Age > 55		-0.044* (0.025)		0.050*** (0.015)		0.005 (0.020)
Income		-0.019** (0.009)		-0.006 (0.005)		-0.025*** (0.006)
Mortgager		0.058*** (0.022)		-0.056*** (0.013)		0.002 (0.018)
Homeowner		0.002 (0.024)		-0.012 (0.015)		-0.011 (0.018)
Income expectations		0.010 (0.017)		0.010 (0.010)		0.020 (0.014)
Constant	0.636*** (0.018)	0.811*** (0.099)	0.146*** (0.009)	0.218*** (0.062)	0.782*** (0.014)	1.029*** (0.069)
R-squared	0.03	0.04	0.02	0.04	0.02	0.04
Observations	3444	3370	3444	3370	3444	3370

Notes: Standard errors in parentheses. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$. Observations are weighted using survey weights.

to directly contrast the MPC out of the hypothetical income gain with the MPC out of an actual income gain. Figure B7 shows that the distribution of MPCs out of tax refunds and hypothetical income gains are similar. The average MPC out of tax refunds is slightly higher because more respondents indicate that they would spend the entire refund. Given that tax refunds are usually significantly lower than 10% of annual income and MPCs out of gains tend to be negatively correlated with the size of the transfer, this is not surprising.

Financial literacy: One might wonder if households are sufficiently financially literate to accurately predict their consumption response to an income change. For this reason, I construct a measure of financial literacy based on seven questions in the SCE that ask respondents to perform simple quantitative exercises. Restricting the sample to only the most financially literate households, Figure B8 shows that the MPC asymmetry is still sizeable.

Intentions vs Actions: Finally, I can directly study to what degree households intended spending coincides with their actual spending. Exploiting the panel dimension of the SCE, I compare ex-ante expenditure plans across seven categories of goods with ex-post purchases four months later. In particular, households are asked to provide an estimate of how likely it is that they will purchase a given good over the next four months. Table B1 shows estimates of a linear probability model and a logit model. There is substantial variation across goods categories, but planned expenditure is a strong predictor of actual expenditure.

4 A model of mental accounts

The data suggest that asymmetric consumption responses are prevalent across most households, irrespective of their wealth levels. These asymmetries allude to a systematic behavioural pattern that induces households to save large fractions of income windfalls, but deters households from using savings to buffer income losses. A conventional consumption framework would predict the former, but not the latter. Instead, I propose an extension of the standard consumption model that incorporates mental accounting (Shefrin and Thaler, 1988; Thaler, 1990; McDowall, 2019).⁸

4.1 General framework

Households hold different mental accounts with regards to their current asset and current income position. Funds pertaining to the mental account for assets are not perfect substitutes for funds in the mental account for income. The partition between the income and the asset account is given by a savings rule. This rule can be thought of as a mental rule-of-thumb that households

⁸Appendix E discusses why several extensions of the standard framework and common behavioural theories fail to generate quantitatively large MPC asymmetries for households that are not liquidity-constrained.

use to facilitate decision-making. Negative deviations from the savings rule, i.e. consuming out of the mental account for assets, are assumed to be costly. Hence, households are *dissaving-averse* (DA). The asymmetric cost originates from the idea that mental accounting introduces an explicit ordering of mental accounts in which consuming out of the income account is preferred to consuming out of the asset account.

Formally, I introduce the mental accounting friction through a modified utility function, as in McDowall (2019):⁹

$$u^{DA}(c) = u(c) - \lambda(a)d(a', a^{plan}) \quad (1)$$

$$d(a', a^{plan}) = \begin{cases} 0 & \text{if } a' \geq a^{plan} \\ u(c) - u(c^{plan}) & \text{if } a' < a^{plan} \end{cases} \quad (2)$$

where $u(c)$ denotes a standard utility function over consumption and $\lambda(a)d(a', a^{plan})$ denotes the disutility from deviating from the savings rule. The disutility term consists of two elements: a penalty function $d(a', a^{plan})$, which depends on the deviation of the actual savings decision a' from the savings rule a^{plan} . This can equivalently be re-mapped into the deviation of actual consumption c from the consumption level that the household obtains following strictly the savings plan, c^{plan} . $d(a', a^{plan})$ is specified in such a way that only negative deviations from the savings plan, i.e. dissaving, are penalized. Saving more than planned does not affect household's utility directly. The second element of the disutility term, $\lambda(a)$, denotes the strength of the dissaving-aversion motive. This formulation of the disutility term is convenient as bounding $\lambda \in [0, 1]$ leads to two extreme types of consumption behaviour at each bound, permanent income consumers for $\lambda = 0$ and hand-to-mouth consumers for $\lambda = 1$. The parameter λ is allowed to depend on the asset position to leave open the possibility that mental accounting frictions may vary with wealth.

Discussion: There is ample empirical evidence that households behave as if they were subject to mental accounting. As an early example, Shefrin and Thaler (1988) tests and confirms several predictions of mental accounting for both liquidity-constrained and unconstrained households. A set of more recent studies finds numerous empirical patterns that can be reconciled by mental accounting behaviour, such as low MPCs out of wealth (Chodorow-Reich et al., 2021; Christelis et al., 2021), differences in spending propensities out of capital gains and dividends (Di Maggio et al., 2020), little reaction to income news (Ganong and Noel, 2019; McDowall, 2019; Baugh et al., 2021; Fuster et al., 2021), anticipation-dependence (Thakral and Tô, 2021) and saving rates that are increasing in permanent income (Straub, 2019). Some of these predictions are

⁹While I use a mental accounting model to study MPC asymmetries, McDowall (2019) uses it to explain the timing of consumption responses. My specification differs in the design of the savings rule and allows for a more flexible penalty term, as discussed later.

Table 3: Share of households with budget or savings/debt repayment plan

Percentile of net liquid wealth distribution	0-20	20-40	40-60	60-80	80+
Keeps budget	68.5%	66.3%	70.8%	65.8%	59.8%
Has savings/debt repayment plan	68.9%	66.5%	59.5%	64.5%	53.5%

Notes: Households are coded as keeping a budget if they answer the following question with yes: "Do you have a (family) budget, or otherwise plan your monthly spending and saving?" Households are coded as having a savings/debt repayment plan if they answer either of the following questions with yes: "People budget in different ways. Do you (and your family) generally try to focus more on trying to save regular amounts of money?" or "People budget in different ways. Do you (and your family) generally try to pay off regular amounts of debt?"

discussed in later sections.

Data from the SCE also support the idea that households follow a savings rule. The survey asks participants whether they keep a budget or follow a savings or debt repayment plan. Table 3 shows that a large majority of households indeed keeps a budget and plans their savings and debt repayments, and that this share is decreasing in wealth.

Finally, households that explicitly state to keep a budget or follow a savings or debt repayment plan are associated with a substantially larger MPC asymmetry (Table 4). This suggests that households subject to mental accounting distinguish more strongly between positive and negative income changes in their consumption response.

4.2 Two-period model

This section presents a simple two-period model to illustrate the mechanism through which mental accounting generates asymmetric MPCs. The two-period model also allows me to derive several analytical propositions with regards to the predictions of the model.

Households in this economy live for two periods $t \in \{0, 1\}$ and are born with zero initial wealth.¹⁰ In the first period, households receive income y_0 and decide how much to consume and how much to save. In the second period, households consume their savings from the first period. Households follow a savings plan which is formed endogenously. Utility is logarithmic, augmented by the penalty term for deviating from the savings plan.¹¹ For simplicity, I assume that households face no penalty in the second period and that the dissaving-aversion parameter λ is constant.¹² Finally, households discount the future with the subjective discount factor β

¹⁰This assumption is purely made for expositional clarity. Appendix C.2 solves for the case with initial wealth.

¹¹Appendix C.3 solves for the case with CRRA utility.

¹²The assumption of no dissaving-aversion in the second period is made purely for clarity. Given that it is always optimal to consume all savings in $t = 1$ and that households do not expect any deviations from their savings plan, the results are identical. See also Appendix C.4.

Table 4: MPCs and savings plans

	(1) MPC^{Asy}	(2) MPC^+	(3) MPC^-
Keeps budget	0.093*** (0.02)	-0.045*** (0.01)	0.049*** (0.01)
Has savings/debt repayment plan only	0.049** (0.02)	-0.043*** (0.01)	0.006 (0.02)
R-squared	0.06	0.05	0.04
Observations	3370	3370	3370
Controls	YES	YES	YES

Notes: Households are coded as keeping a budget if they answer the following question with yes: "Do you have a (family) budget, or otherwise plan your monthly spending and saving?" Households are coded as having a savings/debt repayment plan only if they answer either of the following questions with yes: "People budget in different ways. Do you (and your family) generally try to focus more on trying to save regular amounts of money?" or "People budget in different ways. Do you (and your family) generally try to pay off regular amounts of debt?" but not: "People budget in different ways. Do you (and your family) generally try to pay off regular amounts of debt?". Controls include net liquid wealth, income, housing status, age and income expectations. Standard errors in parentheses. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$. Observations are weighted using survey weights.

and save at an exogenous gross interest rate R . This yields the following problem:

$$\max_{c_0, c_1, a_0^{plan}} u(c_0) - \lambda d(a_0, a_0^{plan}) + \beta u(c_1) \quad (3)$$

$$\text{s.t. } c_0 + a_0 = y_0; \quad c_1 = Ra_0; \quad (4)$$

Taking derivatives with respect to a_0 yields the following Euler equation:

$$\beta R u'(c_1) = \begin{cases} u'(c_0) & \text{if } a_0 \geq a_0^{plan} \\ (1 - \lambda)u'(c_0) & \text{if } a_0 < a_0^{plan} \end{cases} \quad (5)$$

The savings decision governs which Euler equation the household faces. Saving less than planned reduces marginal utility today by a factor $1 - \lambda$. Saving weakly more than planned preserves the standard Euler equation. Combining Euler equation and budget constraint, we can derive an expression for c_0 :

$$c_0 = \begin{cases} \frac{y_0}{1+\beta} & \text{if } a_0 \geq a_0^{plan} \\ \frac{y_0}{1+\frac{\beta}{1-\lambda}} & \text{if } a_0 < a_0^{plan} \end{cases} \quad (6)$$

The final element that is missing is the savings plan itself. I assume that the household's savings plan is given by the optimal savings decision in an equivalent problem without mental

accounting. That is, the savings plan is formed endogenously based on the household's current income and wealth position.¹³ Formally, it is derived as the solution to the following problem:

$$\max_{c_0, c_1} u(c_0) + \beta u(c_1) \quad (7)$$

$$\text{s.t. } c_0 + a_0 = y_0; \quad c_1 = Ra_0; \quad (8)$$

which yields an optimal savings allocation in period 0, a_0^* :

$$a_0^* = y_0 - \frac{y_0}{1 + \beta} \equiv a_0^{plan} \quad (9)$$

With the definition of the savings plan at hand, we can define planned consumption as the level of consumption if one strictly followed the savings plan:

$$c_0^{plan} = y_0 - a_0^{plan} \quad (10)$$

This definition will be useful for providing the intuition behind the results. To derive an expression for the MPC, assume that income in $t = 0$ unexpectedly changes by a fraction ϵ . Additionally, assume that households classify the income shock as a change in their mental account for income but not in their mental account for assets. Formally, this implies that the savings rule does not change in response to the income change. Then, we can compute the MPC as:

$$MPC = \frac{\Delta c_0}{\Delta y_0} = \frac{\tilde{c}_0(y_0 + \epsilon y_0) - c_0(y_0)}{\epsilon y_0} = \begin{cases} \frac{1}{1+\beta} & \text{if } \epsilon \geq 0 \\ \min \left\{ \frac{1}{1+\beta} \left(\frac{1+\epsilon}{\epsilon} \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} - \frac{1}{\epsilon} \right), 1 \right\} & \text{if } \epsilon < 0 \end{cases} \quad (11)$$

This yields the following proposition:

Proposition 1 (MPC asymmetry): Define MPC^+ as the MPC out of positive income changes ($\epsilon > 0$) and MPC^- as the MPC out of negative income changes ($\epsilon < 0$). Then $MPC^- > MPC^+$ for any level of dissaving-aversion $\lambda \in (0, 1]$ and size of the income change $|\epsilon| \in (0, 1)$.

Proof. See Appendix C.1. □

Proposition 1 suggests that the MPC depends on the direction of the income change, i.e. the sign of ϵ . To gain intuition for this result, suppose y_0 increases by a fraction ϵ . This increase does not move planned savings, but it moves planned consumption by ϵy_0 . The increase in planned consumption relaxes the dissaving constraint and additional consumption is not penalized up

¹³Other savings plans could also generate asymmetric MPCs. This formulation, however, yields the cleanest analytical results.

to an increase of ϵy_0 . Hence, as long as the household does not want to increase consumption beyond the increase in income, we recover the standard MPC without mental accounting. Now suppose y_0 decreases by a fraction ϵ . Again, this decrease does not move planned savings, but it reduces planned consumption by ϵy_0 . Any consumption beyond planned consumption is now penalized by $1 - \lambda$ in terms of marginal utility. This results in an MPC out of losses that is higher than the MPC out of gains.

The magnitude of the asymmetry depends on the degree of dissaving-aversion λ and the size of the shock ϵ . In the extreme case where $\lambda = 0$, we recover the standard model without any asymmetries. With $\lambda = 1$, the MPC out of losses is 1 and the household behaves as hand-to-mouth in response to negative income shocks. The next proposition elaborates on the role of ϵ .

Proposition 2 (MPC and shock size): MPC^- is decreasing in the size of the income shock, $\frac{\partial MPC^-}{\partial \epsilon} > 0$, for any level of dissaving-aversion $\lambda \in (0, 1]$ and income shock $\epsilon \in (-1, 0)$. MPC^+ is independent of the size of the shock.

Proof. See Appendix C.1. □

Proposition 2 states that the MPC out of losses is lower for larger shocks. Given that the MPC out of gains is independent of the shock size, the MPC asymmetry is decreasing in the size of the shock. Intuitively, the larger the income loss, the larger is the reduction in planned consumption and as such, the higher is the marginal utility of consuming beyond planned consumption. This reduces the decrease in consumption following the income loss and therefore decreases the MPC. An alternative interpretation not captured in this simple framework is that larger shocks induce agents to update their savings plans (due to increased salience, for example), which brings the MPC closer to the one in a conventional model. The negative relation between MPC^- and shock size established here goes slightly against conventional predictions of a positive relation. However, these predictions are usually based on theories of liquidity-constraints from which I abstract here.

The next two propositions illustrate that the mental accounting model also makes predictions about the MPC out of wealth and income news that are consistent with the data.

Proposition 3 (MPC out of wealth): The MPC out of wealth, $\frac{\Delta c_0}{\Delta w_0}$, is smaller than the MPC out of income if the change in wealth enters the mental account for assets: $MPC^{+,wealth} < MPC^+$ and $MPC^{-,wealth} < MPC^-$ for any level of dissaving-aversion $\lambda \in (0, 1]$ and income change $|\epsilon| \in (0, 1)$.

Proof. See Appendix C.1. □

Proposition 3 shows that the mental accounting model generates MPCs out of wealth that are smaller than MPCs out of income, for both gains and losses, under the assumption that the wealth shock is classified as a change in the mental account for assets. This prediction is supported by a large body of empirical evidence (Di Maggio et al., 2020; Chodorow-Reich et al., 2021; Christelis et al., 2021). While most theoretical explanations are based on differences in liquidity between income and wealth, the mental accounting model introduces non-fungibility by assigning income and wealth shocks to different mental accounts, i.e. through differential responses of the savings plan. Because unexpected changes in wealth shift the savings plan one-to-one with wealth, and therefore leave planned consumption unchanged, consuming out of additional wealth is penalized and yields lower MPCs out of wealth compared to income. Similarly, because planned consumption is unchanged, wealth losses require lower reductions in consumption compared to income losses and yield again lower MPCs.

Proposition 4 (MPC out of income news): The MPC out of income news, $\frac{\Delta c_0}{\Delta y_1}$, is lower than the MPC out of contemporary income changes if income news enter the mental account for future income: $MPC^{+,news} < MPC^+$ and $MPC^{-,news} < MPC^-$ for any level of dissaving-aversion $\lambda \in (0, 1]$, income change $|\epsilon| \in (0, 1)$ and gross interest rate $R \geq 1$.

Proof. See Appendix C.1. □

Proposition 4 suggests that also MPCs out of income news are smaller than MPCs out of contemporaneous income changes under the assumption that news about future income are classified as a change in the mental account for future income. Several papers have documented the insensitivity of consumption to income news (Ganong and Noel, 2019; McDowall, 2019; Baugh et al., 2021; Fuster et al., 2021). Conventional consumption models can explain lower MPCs out of positive income news for liquidity-constrained households, but not for unconstrained households and neither for constrained nor unconstrained households in the case of negative income news.¹⁴ The mental accounting model in contrast generates low MPCs out of news for all households without relying on liquidity-constraints. The non-fungibility between current and future income is again introduced through assignment to different mental accounts, i.e. the response of the savings plan: the savings plan does not change in response to news about future income, similarly to the response to changes in current income. Because the income change only materializes in the next period, planned consumption does not change either. Therefore, any increase in consumption is penalized yielding lower MPCs out of positive news than out of current income gains. Similarly, maintaining the current consumption level is not penalized in response to negative income news compared to current income losses, yielding lower MPCs out of negative news.

¹⁴Except for minor differences in MPCs out of current income and news due to discounting of future income receipts.

5 Quantitative model

This section introduces a quantitative life-cycle model with mental accounting that replicates several empirical MPC patterns and studies implications for fiscal policy.

5.1 Life-cycle model with mental accounting

To explore to what extent the mental accounting model can quantitatively match the MPC asymmetry observed in the data, I incorporate mental accounting into a life-cycle model with idiosyncratic income risk and borrowing constraints.

Model environment: Time is discrete. The economy is populated by a continuum of households, indexed by i . Households live for J periods and work for JR periods after which they retire. While working, households receive a stochastic income $y_{i,t}$. Households can save in a risk-free asset a that pays an interest rate r . Borrowing is not allowed in this economy, i.e. $\underline{a} = 0$. Households have mental accounting preferences given by:

$$u^{DA}(c_{i,t}) = u(c_{i,t}) - \lambda_0 e^{a_{i,t} \lambda_1} d(a_{i,t+1}, a_{i,t+1}^{plan}) \quad (12)$$

In contrast to the two-period model, I allow the strength of the dissaving-aversion parameter to vary with the level of wealth that the household holds. In particular, dissaving-aversion is modelled as an exponential function with level parameter λ_0 and decay parameter λ_1 . This allows the model to flexibly capture two potential features: the covariance between wealth and dissaving-aversion at the intensive margin, i.e. the same household exhibiting different degrees of dissaving-aversion for different levels of wealth or, at the extensive margin, different shares of behavioural households compared to fully rational households at different levels of wealth.

Log income is given by a deterministic life-cycle component \bar{y} and a stochastic component that is modelled as a persistent-transitory process, where the persistent component follows an AR-1 process. The innovations to the persistent and transitory component are orthogonal to each other and independent over time and across households.

$$\log Y_{i,t} = \bar{y}_t + z_{i,t} + e_{i,t} \quad (13)$$

$$z_{i,t} = \rho_z z_{i,t-1} + u_{i,t}, \quad u_{i,t} \sim N(0, \sigma_z^2) \quad e_{i,t} \sim N(0, \sigma_e^2) \quad (14)$$

This yields the following recursive formulation:

$$V(j, z, e, a) = \max_c u^{DA}(c) + \beta EV(j+1, z', e', a') \quad (15)$$

$$\text{s.t.} \quad c + a' = (1+r)a + y, \quad a' \geq \underline{a} \quad (16)$$

where j denotes age. The final element that needs to be defined is the mental partition between income and assets, i.e. the savings plan. Similarly to the stylized framework in Section 4, I specify the savings plan as the optimal savings policy from an equivalent household problem as in Equation 15 in which the transitory shock is set to zero: the savings plan responds to persistent, but not to transitory income changes.

$$a_{i,t}^{plan} = \tilde{a}^*(j, z, e = 0, a) \quad \text{from} \quad \tilde{V}(j, z, e, a) = \max_c u(c) + \beta \mathbb{E} \tilde{V}(j+1, z', e', a') \quad \text{s.t.} \quad (16)$$

Under this assumption, transitory changes to income are mentally classified as income. Persistent changes, in contrast, are partly assigned to the mental account for income and partly to the mental account for assets. The extent of this partition depends on the position in the life-cycle and wealth. Intuitively, this specification of the savings rule provides households with the flexibility to update their mental accounts in response to important events, but preserves heuristic thinking in less impactful situations.

Calibration: Table 5 provides an overview of the calibrated parameters. I first calibrate several parameters outside of the model. The model period is one year. Households work for 40 years and then spend 20 years in retirement. The interest rate is set to 0.02. The degree of risk aversion γ is set to 2. The deterministic income component is estimated from PSID data by regressing the logarithm of income on a cubic polynomial in age and time dummies to control for trends in income over time. The persistence and variance of the stochastic processes are taken from Kaplan and Violante (2022). Retirement income depends on the employment history of households. It is determined by the persistent component of income earned in the final period before retirement and fluctuates with the transitory income state. The replacement rate is set to 0.6. Population shares are calibrated to match the age distribution in the SCE sample. Households' initial asset holdings are chosen to approximate the net asset holdings of households in the SCE between ages 25-30.

The parameters β , λ_0 and λ_1 are calibrated using simulated method of moments. I set the discount factor β to match the average net wealth-to-income ratio in the SCE sample.¹⁵ Disciplining the dissaving-aversion parameters is more intricate. I calibrate the level and decay parameter to match two moments of the data: the average MPC out of income losses and the ratio of households that follow a savings plan between the bottom and top quintile of the wealth distribution from Table 3. The latter moment aims to capture the degree of behavioural frictions across the wealth distribution. It allows me to externally discipline the gradient of the dissaving-motive without any targeting of MPCs. The moment selection is conservative in the sense that I target neither the MPC out of gains, the asymmetry of MPCs nor the behaviour of MPCs across the wealth distribution.

¹⁵Calibrating the model to net liquid wealth instead of net total wealth yields qualitatively similar results.

Table 5: Calibration

Parameter	Description	Value	Source/Target
<i>External</i>			
γ	Risk aversion	2	Standard
J	Length of life-cycle	60	Standard
JR	Length of working-life	40	Standard
\bar{y}	Life-cycle income profile	Cubic polynomial	PSID
ω	Replacement rate	0.6	Standard
r	Interest rate	0.02	Standard
ρ_z	Persistence of z_t	0.953	PSID (Kaplan and Violante, 2022)
σ_z^2	Variance of z_t	0.0422	PSID (Kaplan and Violante, 2022)
σ_e^2	Variance of e_t	0.0494	PSID (Kaplan and Violante, 2022)
\underline{a}	Borrowing limit	0	Standard
<i>Internal</i>			
β	Discount factor	0.93	Avg. net wealth-to-income
λ_0	Dissaving aversion - level	0.69	Avg. MPC^-
λ_1	Dissaving aversion - decay	-0.015	Top-bottom ratio of households with savings plan

The level of dissaving-aversion λ_0 is calibrated to 0.69, with a moderate decay in wealth of $\lambda_1 = -0.015$. The average level of dissaving-aversion across the simulated households of $\bar{\lambda} = 0.63$ turns out to be above the one estimated in McDowall (2019), who finds a value of 0.346. This is not surprising given that the models are not only structurally different across several dimensions - specification of the savings rule, the dissaving-aversion motive and other life-cycle components - but are also calibrated based on different moments. Targeting a lower MPC level as in McDowall (2019) would result in a lower dissaving-aversion parameter. Table D1 shows that the model moments match the targeted data moments very well.

5.2 Results

This section presents the main predictions of the quantitative model. It shows that the model addresses a set of empirical puzzles that go beyond the predictions of the stylized model in Section 4.

MPC asymmetry: I compute MPCs out of income gains and losses by simulating households' consumption paths in response to an exogenous increase and decrease in income by 10 percent, in line with the survey questions from the empirical section. Figure 5 plots the model-generated MPCs across the wealth distribution. The model matches the empirically observed MPCs closely. In particular, it successfully generates large MPCs out of income losses that are decreasing moderately in wealth. At the same time, it produces consumption responses to income gains that are substantially smaller. Figure 6 makes this explicit by plotting the asymmetry itself.

Comparing MPCs from the mental accounting model to MPCs from a model without be-

havioural frictions illustrates the importance of including mental accounting preferences.¹⁶ The standard model generates much smaller MPCs out of income losses than the data suggest, in particular for unconstrained households. It furthermore fails to generate sizeable differences between MPCs out of gains and out of losses. Hence, it neither matches the asymmetry for liquidity-constrained households nor any of the empirically observed MPCs for unconstrained households.


The mental accounting model also matches fairly well the distribution of MPCs across households. Figure 7 shows that the model produces MPCs out of gains and losses that are concentrated around 0 and 1 respectively, closely resembling the patterns in the data. Figure 8 furthermore shows that the model not only replicates the asymmetry in the cross-section but also at the household level. It predicts a high share of households with fully asymmetric MPCs, i.e. asymmetries of close to one. At the same time, it also predicts a significant share of households with symmetric MPCs, i.e. households resembling permanent income consumers.

MPC out of gains: Two observations regarding the MPC out of income gains deserve further discussion. First, the mental accounting model not only generates higher MPCs out of losses compared to the standard model, but also higher MPCs out of gains across all levels of wealth. This is due to some households saving less than prescribed by their savings plan. These households would ideally consume more, but do not do so because they are dissaving-averse. Being constrained by their savings plan, they behave similarly to liquidity-constrained households - any additional income windfall is largely consumed. Hence, mental accounting provides a potential rationale for the empirically observed large consumption response of households with high levels of liquid wealth. Second, the model-generated MPCs out of gains for constrained households are higher than what the SCE data suggest, but largely in line with findings in other empirical studies (Parker et al., 2013). This is due to the mental accounting model preserving the negative relation between MPCs out of gains and wealth induced by borrowing constraints.

MPC and shock size: Figure 9 shows that the size of the income shock matters for both the level of the MPC and the relation with wealth. It compares consumption responses to a 3% and 30% income change, i.e. a relative small and relatively large shock compared to the baseline scenario. Larger shocks generate larger asymmetries in the mental accounting framework, driven by a larger consumption response to income losses. This helps reconcile why empirical estimates of MPC asymmetries differ across studies - differently sized shocks generate different degrees of asymmetries. Fuster et al. (2021) for example finds an average asymmetry of 0.25, which is around half the size of the asymmetry found in this paper. At the same time, they study an income shock that is substantially smaller - 500\$ compared to 10% of annual income.

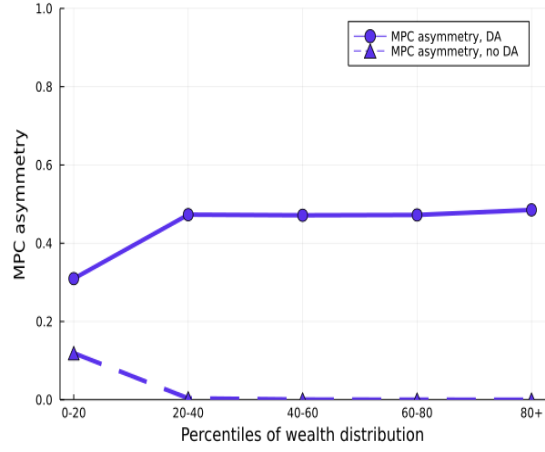
¹⁶Because mental accounting preferences slightly change the distribution of assets in the economy, I recalibrate the discount factor in the model without mental accounting preferences to match the same average net wealth-to-income ratio.

Figure 5: Model MPCs




Notes: MPCs are computed by simulating a transitory 10% income shock. DA refers to the model with mental accounting preferences while no DA refers to the frictionless model.

Figure 6: Model MPC asymmetries



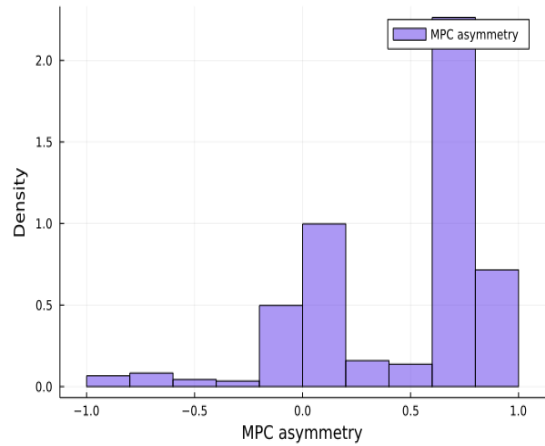
Notes: MPCs are computed by simulating a transitory 10% income shock. DA refers to the model with mental accounting preferences while no DA refers to the frictionless model.

Figure 7: Distribution of model MPCs



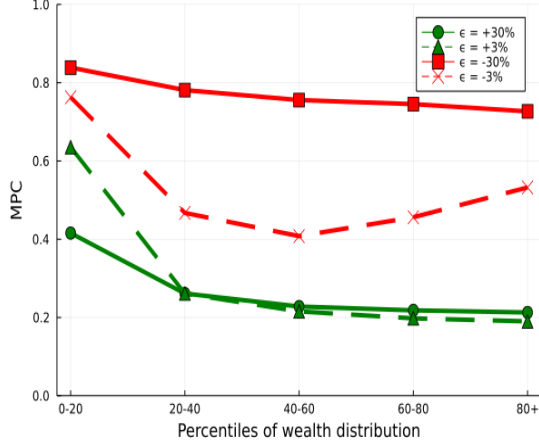
Notes: MPCs are computed by simulating a transitory 10% income shock. Due to numerical error in the simulation of MPCs, some households have MPCs that are marginally below 0 or above 1. This are recoded to 0 and 1, respectively.

Figure 8: Distribution of model MPC asymmetries




Notes: MPCs are computed by simulating a transitory 10% income shock. Due to numerical error in the simulation of MPCs, some households have MPCs that are marginally below 0 or above 1. This are recoded to 0 and 1, respectively.

Figure 9: Model MPC by size of income shock



Notes: MPCs are computed by simulating transitory proportional income shocks of different sizes.

Figure 10: Model MPC out of wealth and income news



Notes: MPCs are computed by simulating (i) a transitory 10% shock to current income, (ii) news about a 10% shock to income next period and (iii) a 10% shock to current wealth.

How do we reconcile these findings with Proposition 2 that suggests that MPCs out of losses are decreasing in the size of the shock? In the stylized framework, the household is exactly at the dissaving-constraint, i.e. actual savings equal planned savings before the shock hits. In the quantitative model, however, some households are above and some are below the constraint. A larger shock pushes a larger fraction of initially unconstrained households into the constrained region and increases their MPC. The stronger compositional effect induced by larger shocks more than off-sets the decrease in MPCs for households that are already constrained and thus generates MPCs out of losses that are larger on average.

Finally, larger shocks flatten the relation between MPCs out of gains and wealth. This is particularly relevant in light of recent evidence by Andreolli and Surico (2021) that finds a negative relation between MPCs out of gains and wealth for small shocks, but a flat relation for large shocks. They interpret this finding through a model with non-homothetic preferences, but the mental accounting model seems to, at least qualitatively, produce similar patterns.

MPC out of wealth: Consistent with the stylized theoretical framework, the quantitative model generates relatively small MPCs out of wealth. Figure 10 shows that the average MPC out of both wealth gains and losses is several times smaller than the respective MPC out of income. The model without mental accounting, instead, fails to generate differences between MPCs out of income and wealth.¹⁷ Because wealth and income are perfectly fungible in the

¹⁷Introducing transaction costs to the standard model can generate lower MPCs out of wealth for illiquid wealth, but not liquid wealth. The mental accounting model, however, predicts lower MPCs for both types of

standard framework, the MPC out of wealth is identical to the MPC out of income.

MPC out of news: The mental accounting model also generates MPCs out of income news that are substantially smaller than MPCs out of current income (Figure 10).¹⁸ Moreover, the model preserves the sign asymmetry for MPCs out of news. This resonates with the empirical evidence in Fuster et al. (2021) that finds that households adjust consumption more in response to negative than to positive income news.

The lower sensitivity of consumption to income news also relates to recent empirical evidence on the reaction to predictable income changes. McDowall (2019) finds that consumption primarily responds upon the receipt of income, but not upon the arrival of income news and explains this fact with a model of mental accounting. Ganong and Noel (2019) finds similar evidence for predictable income losses, but explains the findings with present bias. Present bias, in contrast to mental accounting, cannot explain the sluggish reaction of unconstrained households to positive income news, however. Baugh et al. (2021) studies asymmetric responses to positive and negative predictable income changes and finds the reverse asymmetry that I find for unpredictable income changes: consumption responds more to predictable gains than to losses. The stronger initial response of consumption to negative compared to positive income news could explain why upon receipt of income the response to positive shocks is larger than to negative ones.

The model without mental accounting fails to generate substantial differences between MPCs out of current income and income news for unconstrained households. The MPC out of income news is identical to the MPC out of income adjusted for discounting, and as such only marginally lower.

Consumption-savings dynamics: Introducing mental accounting preferences to match consumption responses does not come at the cost of missing other moments. The consumption-savings dynamics are close to the ones of the frictionless life-cycle model. Figures 11 and 12 show that the mental accounting model produces a comparable dispersion of wealth and consumption, while Figure 13 shows that it also predicts a similar life-cycle profile of consumption and savings.


5.3 Implications for fiscal policy

Large asymmetries in MPCs have important implications for the design of fiscal policies, in particular for redistributive measures. The commonly held view that redistribution from the rich to the poor boosts aggregate demand through reallocation of resources from low to high MPC

wealth.

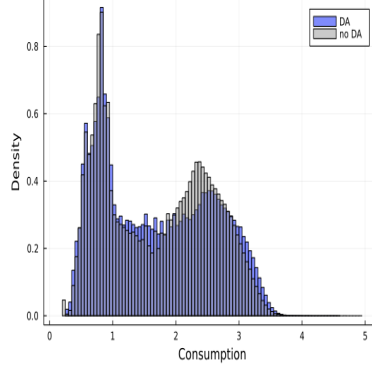
¹⁸The MPC out of income news is defined as the contemporaneous consumption response to news about an income shock one period ahead.

Figure 11: Wealth dispersion



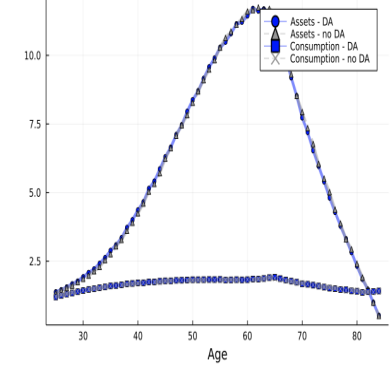
Notes: DA refers to the model with mental accounting preferences while no DA refers to the frictionless model.

Figure 12: Consumption dispersion



Notes: DA refers to the model with mental accounting preferences while no DA refers to the frictionless model.

Figure 13: Consumption and savings over the life-cycle



Notes: DA refers to the model with mental accounting preferences while no DA refers to the frictionless model.

households does not necessarily apply. If rich households have large MPCs out of income losses, their reduction in consumption could more than compensate for the increase in consumption by the poor. To assess the quantitative significance of this argument, this section compares the effectiveness of two different fiscal policies.

Policy design: I evaluate a policy in which the government sends targeted lump-sum transfers to the bottom half of the income distribution. The size of the transfer is calibrated to roughly match the stimulus checks that were disbursed as part of the COVID-19 Stimulus Package in the US.¹⁹ The transfers are financed in two different ways: a one-off proportional income tax on the top 25 percent of the income distribution and a one-off proportional wealth tax on the top 25 percent of the wealth distribution.

Results: Table 6 reports the percentage change in aggregate consumption following the introduction of the policy for each type of financing scheme and compares it to the change in aggregate consumption in a model without mental accounting preferences. It also reports the tax rate that is required to finance the transfers.

The first policy design illustrates that redistributive measures can be less effective when MPCs are asymmetric. A policy in which transfers are financed through an income tax on the top 25% of the income distribution *lowers* aggregate consumption by 0.01 percent in the model with mental accounting preferences. It is substantially less effective than in the standard model

¹⁹Eligible individuals received a payment of \$1,400 (\$2,800 for married couples), plus an additional \$1,400 per eligible child. With an average household size in the US of around 2.5, this results in a payment of roughly \$3,500 per household. This is around five percent of median income in the SCE sample, and as such roughly half the size of the hypothetical 10 percent shock for the median household.

Table 6: Effects of redistributive fiscal policy

	% change in aggregate consumption	Required tax rate
<i>Income tax on the rich</i>		
DA model	-0.01	3.7
Standard model	0.25	3.7
<i>Wealth tax on the wealthy</i>		
DA model	0.57	0.7
Standard model	0.24	0.7

Notes: This table reports changes in aggregate consumption following lump-sum transfers to the bottom 50 percent of the income distribution financed by a (i) one-off proportional income tax on the top 25 percent of the income distribution, and (ii) a one-off proportional wealth tax on the top 25 percent of the wealth distribution.

due to larger consumption reductions by the rich.

The second policy design suggests that the type of tax matters. Financing the transfers through a wealth tax instead of an income tax leads to a large increase in aggregate consumption as consumption is much less sensitive to changes in wealth than to changes in income. The effect is also stronger than in the standard model because MPCs out of income gains are larger.

6 Conclusion

This paper documented that consumption responds asymmetrically to changes in income. Consumption is smoothed less in response to income losses than in response to income gains. A simple extension of the standard consumption framework that incorporates mental accounting can generate the empirically observed asymmetry. The model also predicts consumption responses to changes in wealth and income news that are consistent with empirical evidence.

A quantitative evaluation of the model with mental accounting illustrated that the implications for fiscal policy are far-reaching. Redistributive fiscal policy can be less effective in stimulating aggregate demand than in a conventional framework if the consumption of households that are taxed to finance the policy is sensitive to income changes. As such, adequate targeting of population segments is critical both on the spending and the financing side for effective fiscal policy. More broadly, asymmetric MPCs also suggest that fiscal contractions could translate into stronger aggregate consumption responses than fiscal expansions, as for example documented in [Barnichon et al. \(2022\)](#).

References

- Andreolli, M. and Surico, P. (2021). Less is more: Consumer spending and the size of economic stimulus payments.
- Arellano, M., Blundell, R., and Bonhomme, S. (2017). Earnings and consumption dynamics: a nonlinear panel data framework. *Econometrica*, 85(3):693–734.
- Attanasio, O., Kovacs, A., and Moran, P. (2020). Temptation and commitment: Understanding hand-to-mouth behavior. Technical report, National Bureau of Economic Research.
- Barnichon, R., Debortoli, D., and Matthes, C. (2022). Understanding the size of the government spending multiplier: It’s in the sign. *The Review of Economic Studies*, 89(1):87–117.
- Baugh, B., Ben-David, I., Park, H., and Parker, J. A. (2021). Asymmetric consumption smoothing. *American Economic Review*, 111(1):192–230.
- Blundell, R., Pistaferri, L., and Preston, I. (2008). Consumption inequality and partial insurance. *American Economic Review*, 98(5):1887–1921.
- Boutros, M. (2019). Household finances and fiscal stimulus in 2008. *Available at SSRN 3716571*.
- Boutros, M. (2022). Windfall income shocks with finite planning horizons.
- Bracha, A. and Cooper, D. (2014). Asymmetric responses to income changes: The payroll tax increase versus tax refund in 2013. *Economics Letters*, 124(3):534–538.
- Bunn, P., Le Roux, J., Reinold, K., and Surico, P. (2018). The consumption response to positive and negative income shocks. *Journal of Monetary Economics*, 96:1–15.
- Chodorow-Reich, G., Nenov, P. T., and Simsek, A. (2021). Stock market wealth and the real economy: A local labor market approach. *American Economic Review*, 111(5):1613–57.
- Christelis, D., Georgarakos, D., Jappelli, T., Pistaferri, L., and Van Rooij, M. (2019). Asymmetric consumption effects of transitory income shocks. *The Economic Journal*, 129(622):2322–2341.
- Christelis, D., Georgarakos, D., Jappelli, T., Pistaferri, L., and Van Rooij, M. (2021). Heterogeneous wealth effects. *European Economic Review*, 137:103805.
- Commault, J. (2020). Does consumption respond to transitory shocks? *American Economic Journal: Macroeconomics*.
- Di Maggio, M., Kermani, A., and Majlesi, K. (2020). Stock market returns and consumption. *The Journal of Finance*, 75(6):3175–3219.

- Fagereng, A., Holm, M. B., and Natvik, G. J. (2021). Mpc heterogeneity and household balance sheets. *American Economic Journal: Macroeconomics*, 13(4):1–54.
- Fuster, A., Kaplan, G., and Zafar, B. (2021). What would you do with \$500? spending responses to gains, losses, news, and loans. *The Review of Economic Studies*, 88(4):1760–1795.
- Gabaix, X. (2014). A sparsity-based model of bounded rationality. *The Quarterly Journal of Economics*, 129(4):1661–1710.
- Ganong, P. and Noel, P. (2019). Consumer spending during unemployment: Positive and normative implications. *American Economic Review*, 109(7):2383–2424.
- Graziani, G., Van Der Klaauw, W., and Zafar, B. (2016). Workers’ spending response to the 2011 payroll tax cuts. *American Economic Journal: Economic Policy*, 8(4):124–59.
- Gul, F. and Pesendorfer, W. (2001). Temptation and self-control. *Econometrica*, 69(6):1403–1435.
- Güvenen, F., Karahan, F., Ozkan, S., and Song, J. (2021). What do data on millions of us workers reveal about lifecycle earnings dynamics? *Econometrica*, 89(5):2303–2339.
- Ilut, C. L. and Valchev, R. (2020). Economic agents as imperfect problem solvers. Technical report, National Bureau of Economic Research.
- Jappelli, T. and Pistaferri, L. (2010). The consumption response to income changes. *Annu. Rev. Econ.*, 2(1):479–506.
- Jappelli, T. and Pistaferri, L. (2014). Fiscal policy and mpc heterogeneity. *American Economic Journal: Macroeconomics*, 6(4):107–36.
- Kaplan, G. and Violante, G. L. (2010). How much consumption insurance beyond self-insurance? *American Economic Journal: Macroeconomics*, 2(4):53–87.
- Kaplan, G. and Violante, G. L. (2014). A model of the consumption response to fiscal stimulus payments. *Econometrica*, 82(4):1199–1239.
- Kaplan, G. and Violante, G. L. (2022). The marginal propensity to consume in heterogeneous agent models. Technical report, National Bureau of Economic Research.
- Kaplan, G., Violante, G. L., and Weidner, J. (2014). The wealthy hand-to-mouth. Technical report, National Bureau of Economic Research.
- Kőszegi, B. and Rabin, M. (2006). A model of reference-dependent preferences. *The Quarterly Journal of Economics*, 121(4):1133–1165.

- Kőszegi, B. and Rabin, M. (2009). Reference-dependent consumption plans. *American Economic Review*, 99(3):909–36.
- Kueng, L. (2018). Excess sensitivity of high-income consumers. *The Quarterly Journal of Economics*, 133(4):1693–1751.
- Laibson, D., Maxted, P., and Moll, B. (2021). Present bias amplifies the household balance-sheet channels of macroeconomic policy. Technical report, National Bureau of Economic Research.
- Laibson, D., Maxted, P., and Moll, B. (2022). A simple mapping from mpes to mpes. Working Paper 29664, National Bureau of Economic Research.
- Lewis, D. J., Melcangi, D., and Pilossoph, L. (2019). Latent heterogeneity in the marginal propensity to consume. *FRB of New York Staff Report*, (902).
- Lian, C. (2020). Mistakes in future consumption, high mpes now. Technical report, UC Berkeley mimeo.
- McDowall, R. A. (2019). Consumption behavior across the distribution of liquid assets. Technical report, Working paper, NYU.
- Olafsson, A. and Pagel, M. (2018). The liquid hand-to-mouth: Evidence from personal finance management software. *The Review of Financial Studies*, 31(11):4398–4446.
- Pagel, M. (2017). Expectations-based reference-dependent life-cycle consumption. *The Review of Economic Studies*, 84(2):885–934.
- Parker, J. A., Schild, J., Erhard, L., and Johnson, D. S. (2021). Household spending responses to the economic impact payments of 2020: Evidence from the consumer expenditure survey.
- Parker, J. A. and Souleles, N. S. (2019). Reported effects versus revealed-preference estimates: Evidence from the propensity to spend tax rebates. *American Economic Review: Insights*, 1(3):273–90.
- Parker, J. A., Souleles, N. S., Johnson, D. S., and McClelland, R. (2013). Consumer spending and the economic stimulus payments of 2008. *American Economic Review*, 103(6):2530–53.
- Reis, R. (2006). Inattentive consumers. *Journal of Monetary Economics*, 53(8):1761–1800.
- Sahm, C. R., Shapiro, M. D., and Slemrod, J. (2010). Household response to the 2008 tax rebate: Survey evidence and aggregate implications. *Tax Policy and the Economy*, 24(1):69–110.
- Sahm, C. R., Shapiro, M. D., and Slemrod, J. (2015). Balance-sheet households and fiscal stimulus: Lessons from the payroll tax cut and its expiration. Technical report, National Bureau of Economic Research.

- Shapiro, M. D. and Slemrod, J. (2003). Consumer response to tax rebates. *American Economic Review*, 93(1):381–396.
- Shefrin, H. M. and Thaler, R. H. (1988). The behavioral life-cycle hypothesis. *Economic Inquiry*, 26(4):609–643.
- Sims, C. A. (2003). Implications of rational inattention. *Journal of Monetary Economics*, 50(3):665–690.
- Straub, L. (2019). Consumption, savings, and the distribution of permanent income. *Unpublished manuscript, Harvard University*.
- Surico, P. and Trezzi, R. (2019). Consumer spending and property taxes. *Journal of the European Economic Association*, 17(2):606–649.
- Thakral, N. and Tô, L. (2021). Anticipation and consumption. *Available at SSRN 3756188*.
- Thaler, R. H. (1990). Anomalies: Saving, fungibility, and mental accounts. *Journal of Economic Perspectives*, 4(1):193–205.

APPENDIX

A Data

This appendix shows the phrasing of the survey questions and the response options. If households select response option 4-7, they are additionally asked to quantify what percentage they would spend, save or use to pay down debt in case of an income gain, and by how much they would cut spending, savings or increase debt in case of an income loss.

MPC out of income gains: *Suppose next year you were to find your household with 10 percent more income than you currently expect. What would you do with the extra income?*


1. Save or invest all of it
2. Spend or donate all of it
3. Use all of it to pay down debts
4. Spend some and save some
5. Spend some and use part of it to pay down debts
6. Save some and use part of it to pay down debts
7. Spend some, save some, and use some to pay down debts

MPC out of income losses: *Now imagine that next year you were to find yourself with 10% less household income. What would you do?*

1. Cut spending by whole amount
2. Not cut spending at all, but cut my savings by the whole amount
3. Not cut spending, but increase my debt by borrowing the whole amount
4. Cut spending by some and cut savings by some
5. Cut spending by some and increase debt by some
6. Cut savings by some and increase debt by some
7. Cut spending by some, cut savings by some and increase debt by some


B Additional empirical results

Figure B1: MPCs across the net liquid wealth distribution



Notes: MPC asymmetry is defined as the difference between the MPC out of losses and the MPC out of gains. Grey bars indicate 95% confidence bands.

Figure B2: MPC asymmetries across measures of liquidity constraints




Notes: Each line corresponds to the MPC asymmetry across the respective distribution. Grey bars indicate 95% confidence bands. Total net wealth is defined as total assets (liquid assets + retirement funds and housing wealth) - total debt (liquid debt + mortgages). Bank holdings refer to money in checking and savings accounts.

Table B1: Planned vs actual expenditure.

	(1)	(2)	(3)	(4)	(5)	(6)	(7)
	Appliances	Electronics	Furniture	Home repairs	Car	Trips	House
LPM	0.30*** (0.03)	0.37*** (0.02)	0.39*** (0.03)	0.48*** (0.02)	0.41*** (0.03)	0.56*** (0.02)	0.28*** (0.03)
Logit	0.30*** (0.03)	0.37*** (0.02)	0.39*** (0.03)	0.48*** (0.02)	0.41*** (0.03)	0.56*** (0.02)	0.28*** (0.03)
R-squared LPM	0.04	0.05	0.08	0.13	0.07	0.21	0.10
R-squared Logit	0.04	0.04	0.09	0.11	0.07	0.17	0.17
Observations	5704	5693	5683	5691	5673	5690	4741


Notes: Standard errors in parentheses. * $p < 0.1$, ** $p < 0.05$, *** $p < 0.01$.

Figure B3: Marginal propensities to consume and repay debt across debtors and creditors




Notes: Net debtors are defined as households that hold net liquid debt. Percentiles of the wealth distribution are computed conditional on holding positive net liquid wealth. Grey bars indicate 95% confidence bands.

Figure B4: MPC asymmetries across the income distribution




Notes: Grey bars indicate 95% confidence bands.

Figure B5: MPC asymmetries across age groups




Notes: Grey bars indicate 95% confidence bands.

Figure B6: MPC asymmetries across housing status




Notes: Grey bars indicate 95% confidence bands.

Figure B7: MPCs out of tax refunds versus hypothetical scenarios



Notes: Dashed lines denote the average MPC out of gains and tax refunds, respectively.

Figure B8: MPC asymmetries for most financially literate households



Notes: Grey bars indicate 95% confidence bands. Households are coded as financially literate if they answered all questions about financial literacy correctly. This is the case for around one third of the sample (N=1,382)

C Theory

C.1 Proofs

The proofs in this section assume a positive discount factor, i.e. $\beta > 0$.

Derivation of MPCs in the mental accounting model:

The consumption allocation is given by:

$$c_0 = \begin{cases} \frac{y_0}{1+\beta} & \text{if } c_0 \leq c_0^{plan} \\ \frac{y_0}{1+\frac{\beta}{1-\lambda}} & \text{if } c_0 > c_0^{plan} \end{cases}$$

The MPC is given by:

$$MPC = \frac{\Delta c_0}{\Delta y_0} = \frac{\tilde{c}_0(y_0 + \epsilon y_0) - c_0(y_0)}{\epsilon y_0}$$

$\tilde{c}(y_0 + \epsilon y_0)$ denotes the consumption allocation under the savings plan $a_0^{plan}(y_0)$, while $c(y_0 + \epsilon y_0)$ denotes the consumption allocation under the savings plan $a_0^{plan}(y_0 + \epsilon y_0)$. Hence $\tilde{c}(y_0 + \epsilon y_0)$ and $c(y_0 + \epsilon y_0)$ are not necessarily equal.

A positive shock $\epsilon > 0$ increases planned consumption c_0^{plan} by ϵy_0 . Unless the household increases consumption by more than ϵy_0 (which implies a $MPC > 1$) consumption is always weakly below planned consumption. Because it is never optimal to increase consumption by more than ϵy_0 due to consumption smoothing, consumption is indeed always weakly below planned consumption. Hence:

$$MPC^+ = \left(\frac{y_0 + \epsilon y_0}{1 + \beta} - \frac{y_0}{1 + \beta} \right) \frac{1}{\epsilon y_0} = \left(\frac{1 + \epsilon}{1 + \beta} - \frac{1}{1 + \beta} \right) \frac{1}{\epsilon} = \frac{1}{1 + \beta}$$

A negative shock $\epsilon < 0$ decreases planned consumption c_0^{plan} by ϵy_0 . Unless the household decreases consumption by more than ϵy_0 (which implies a $MPC > 1$) consumption is always weakly above planned consumption. Because it is never optimal to decrease consumption by more than ϵy_0 due to consumption smoothing, consumption is indeed always weakly below planned consumption. Hence:

$$\begin{aligned} MPC^- &= \min \left\{ \left(\frac{y_0 + \epsilon y_0}{1 + \frac{\beta}{1-\lambda}} - \frac{y_0}{1 + \beta} \right) \frac{1}{\epsilon y_0}, 1 \right\} = \min \left\{ \left(\frac{1 + \epsilon}{1 + \frac{\beta}{1-\lambda}} - \frac{1}{1 + \beta} \right) \frac{1}{\epsilon}, 1 \right\} \\ &= \min \left\{ \frac{1}{1 + \beta} \left(\frac{1 + \epsilon}{\epsilon} \frac{1 + \beta}{1 + \frac{\beta}{1-\lambda}} - \frac{1}{\epsilon} \right), 1 \right\} \end{aligned}$$

Proposition 1 (MPC asymmetry):

Proof. We want to show that $\min \left\{ \frac{1}{1 + \beta} \left(\frac{1 + \epsilon}{\epsilon} \frac{1 + \beta}{1 + \frac{\beta}{1-\lambda}} - \frac{1}{\epsilon} \right), 1 \right\} > \frac{1}{1 + \beta}$ for $\epsilon \in (-1, 0)$. With regards to the first expression, dividing both sides by $\frac{1}{1 + \beta}$ yields $\frac{1 + \epsilon}{\epsilon} \frac{1 + \beta}{1 + \frac{\beta}{1-\lambda}} - \frac{1}{\epsilon} > 1$. From there, $\frac{1 + \beta}{1 + \frac{\beta}{1-\lambda}} \leq 1$, which is true for any $\lambda \in (0, 1]$. With regards to the second expression, $1 > \frac{1}{1 + \beta}$ for $\beta > 0$. \square

Proposition 2 (Shock size):

Proof. First, we want to show that the derivative of MPC^- with respect to ϵ is increasing for $\epsilon \in (-1, 0)$.

$$\frac{\partial MPC^-}{\partial \epsilon} = \frac{\partial \frac{1}{1 + \beta} \left(\frac{1 + \epsilon}{\epsilon} \frac{1 + \beta}{1 + \frac{\beta}{1-\lambda}} - \frac{1}{\epsilon} \right)}{\partial \epsilon} = \frac{1 + \beta}{1 + \frac{\beta}{1-\lambda}} \frac{\epsilon - (1 + \epsilon)}{\epsilon^2} + \frac{1}{\epsilon^2} = \frac{1}{\epsilon^2} - \frac{1}{\epsilon^2} \underbrace{\frac{1 + \beta}{1 + \frac{\beta}{1-\lambda}}}_{\leq 1} > 0$$

for $\lambda \in (0, 1]$ and $\beta > 0$. Second, $MPC^+ = \frac{1}{1 + \beta}$ and as such does not depend on the income shock ϵ .

□

Proposition 3 (MPC out of wealth):

Proof. Introduce initial wealth w to the problem and assume, for simplicity, $y_0 = 0$. Furthermore, assume that a_0^{plan} changes one-to-one as initial wealth changes. The consumption allocation is then given by:

$$c_0 = \begin{cases} \frac{w}{1+\beta} & \text{if } c_0 \leq c_0^{plan} \\ \frac{w}{1+\frac{\beta}{1-\lambda}} & \text{if } c_0 > c_0^{plan} \end{cases}$$

A positive shock $\epsilon > 0$ increases planned savings a_0^{plan} by ϵw and leaves planned consumption c_0^{plan} unchanged. For any increase in consumption, consumption is therefore always above planned consumption. Furthermore, consumption will never drop in response to a positive income shock due to consumption smoothing. This yields the following MPC^+ out of wealth:

$$\begin{aligned} MPC^{+,wealth} &= \frac{\Delta c_0}{\Delta w} = \max \left\{ \left(\frac{w(1+\epsilon)}{1+\frac{\beta}{1-\lambda}} - \frac{w}{1+\beta} \right) \frac{1}{\epsilon w}, 0 \right\} = \max \left\{ \left(\frac{1+\epsilon}{1+\frac{\beta}{1-\lambda}} - \frac{1}{1+\beta} \right) \frac{1}{\epsilon}, 0 \right\} \\ &= \max \left\{ \frac{1}{1+\beta} \left(\frac{1+\epsilon}{\epsilon} \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} - \frac{1}{\epsilon} \right), 0 \right\} \end{aligned}$$

We want to show that $MPC^{+,wealth} < MPC^+$. With regards to the first expression, we can show that:

$$\frac{1}{1+\beta} \left(\frac{1+\epsilon}{\epsilon} \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} - \frac{1}{\epsilon} \right) < \frac{1}{1+\beta} \rightarrow \frac{1+\epsilon}{\epsilon} \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} - \frac{1}{\epsilon} < 1 \rightarrow \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} < 1$$

for any $\lambda \in (0, 1]$, $\epsilon \in (0, 1)$ and $\beta > 0$. With regards to the second expression, trivially $0 < \frac{1}{1+\beta}$.

A negative shock $\epsilon < 0$ decreases planned savings a_0^{plan} by ϵw and leaves planned consumption c_0^{plan} unchanged. For any decrease in consumption, consumption is therefore always below planned consumption. This yields the following MPC^- out of wealth:

$$MPC^{-,wealth} = \frac{\Delta c_0}{\Delta w} = \left(\frac{w(1+\epsilon)}{1+\beta} - \frac{w}{1+\beta} \right) \frac{1}{\epsilon w} = \frac{1}{1+\beta}$$

We want to show that $MPC^{-,wealth} < MPC^-$.

$$\frac{1}{1+\beta} < \frac{1}{1+\beta} \left(\frac{1+\epsilon}{\epsilon} \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} - \frac{1}{\epsilon} \right)$$

The proof of Proposition 1 shows that this holds for any $\epsilon \in (-1, 0)$ and $\lambda \in (0, 1]$. \square

Proposition 4 (MPC out of income news):

Proof. Introduce income y_1 to the initial problem and assume, for simplicity, $y_0 = 0$. Furthermore, assume that a_0^{plan} does not respond to changes in y_1 . The consumption allocation is then given by:

$$c_0 = \begin{cases} \frac{y_1}{R(1+\beta)} & \text{if } c_0 \leq c_0^{plan} \\ \frac{y_1}{R(1+\frac{\beta}{1-\lambda})} & \text{if } c_0 > c_0^{plan} \end{cases}$$

A positive shock $\epsilon > 0$ to future income y_1 leaves both planned consumption c_0^{plan} and planned savings a_0^{plan} unchanged. For any increase in consumption, consumption is therefore always above planned consumption. Furthermore, consumption will never drop in response to a positive news shock due to consumption smoothing. This yields the following MPC^+ out of income news:

$$\begin{aligned} MPC^{+,news} &= \frac{\Delta c_0}{\Delta y_1} = \max \left\{ \left(\frac{(1+\epsilon)y_1}{R(1+\frac{\beta}{1-\lambda})} - \frac{y_1}{R(1+\beta)} \right) \frac{1}{\epsilon y_1}, 0 \right\} = \max \left\{ \left(\frac{1+\epsilon}{1+\frac{\beta}{1-\lambda}} - \frac{1}{1+\beta} \right) \frac{1}{R\epsilon}, 0 \right\} \\ &= \max \left\{ \frac{1}{1+\beta} \left(\frac{1+\epsilon}{R\epsilon} \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} - \frac{1}{R\epsilon} \right), 0 \right\} \end{aligned}$$

We want to show that $MPC^{+,news} < MPC^+$. With regards to the first expression, we can show that:

$$\frac{1}{1+\beta} \left(\frac{1+\epsilon}{R\epsilon} \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} - \frac{1}{R\epsilon} \right) < \frac{1}{1+\beta} \rightarrow \frac{1+\epsilon}{R\epsilon} \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} - \frac{1}{R\epsilon} < 1 \rightarrow \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} < \frac{1+R\epsilon}{1+\epsilon}$$

for any $\lambda \in (0, 1]$, $\epsilon \in (0, 1)$, $\beta > 0$ and $R \geq 1$. With regards to the second expression, trivially $0 < \frac{1}{1+\beta}$. Following a similar logic, one can also show that $MPC^{+,news} < \frac{1}{R(1+\beta)}$ which is the MPC out of news in a model without dissaving aversion.

A negative shock $\epsilon < 0$ to future income y_1 leaves both planned consumption c_0^{plan} and planned savings a_0^{plan} unchanged. For any decrease in consumption, consumption is therefore always below planned consumption. This yields the following MPC^- out of income news:

$$MPC^{-,news} = \frac{\Delta c_0}{\Delta y_1} = \left(\frac{y_1(1+\epsilon)}{R(1+\beta)} - \frac{y_1}{R(1+\beta)} \right) \frac{1}{\epsilon y_1} = \frac{1}{R(1+\beta)}$$

We want to show that $MPC^{-,news} < MPC^-$:

$$\frac{1}{R(1+\beta)} < \frac{1}{1+\beta} \left(\frac{1+\epsilon}{\epsilon} \frac{1+\beta}{1+\frac{\beta}{1-\lambda}} - \frac{1}{\epsilon} \right)$$

The proof of Proposition 1 shows that this holds for any $\epsilon \in (-1, 0)$, $\lambda \in (0, 1]$ and $R \geq 1$. \square

C.2 DA model with initial wealth

This section studies the relationship between MPCs and wealth in the DA framework. Introducing initial wealth w yields the following problem:

$$\begin{aligned} \max_{c_0, c_1, a_0^{plan}} \quad & u(c_0) - \lambda d(a_0, a_0^{plan}) + \beta u(c_1) \\ \text{s.t.} \quad & c_0 + a_0 = y_0 + w; \quad c_1 = Ra_0 \end{aligned}$$

Solving this problem yields the following consumption allocation: The consumption allocation is given by:

$$c_0 = \begin{cases} \frac{w+y_0}{1+\beta} & \text{if } c_0 \leq c_0^{plan} \\ \frac{w+y_0}{1+\frac{\beta}{1-\lambda}} & \text{if } c_0 > c_0^{plan} \end{cases}$$

The savings plan is derived as in the benchmark problem and given by $a_0^{plan} = w + y_0 - \frac{w+y_0}{1+\beta}$. Following the earlier logic, this yields the following MPCs:

$$MPC^+ = \left(\frac{w + y_0 + \epsilon y_0}{1 + \beta} - \frac{w + y_0}{1 + \beta} \right) \frac{1}{\epsilon y_0} = \frac{1}{1 + \beta}$$

$$MPC^- = \min \left\{ \left(\frac{w + y_0 + \epsilon y_0}{1 + \frac{\beta}{1-\lambda}} - \frac{w + y_0}{1 + \beta} \right) \frac{1}{\epsilon y_0}, 1 \right\}$$

Hence, the MPC^+ does not depend on initial wealth. For MPC^- , take the derivative of the first argument with respect to wealth (with some abuse of notation):

$$\frac{\partial MPC^-}{\partial w} = \left(\frac{w}{1 + \frac{\beta}{1-\lambda}} - \frac{w}{1 + \beta} \right) \frac{1}{\epsilon y_0} > 0$$

for $\epsilon \in (-1, 0)$. Hence, MPC^- is increasing in initial wealth w .

C.3 DA model with CRRA utility

This appendix generalizes the 2-period model to any CRRA utility function $u(c) = \frac{c^{1-\gamma}}{1-\gamma}$. The consumption allocation is given by:

$$c_0 = \begin{cases} \frac{y_0}{1+\beta^{\frac{1}{\gamma}} R^{\frac{1-\gamma}{\gamma}}} & \text{if } c_0 \leq c_0^{plan} \\ \frac{y_0}{1+(\frac{\beta}{1-\lambda})^{\frac{1}{\gamma}} R^{\frac{1-\gamma}{\gamma}}} & \text{if } c_0 > c_0^{plan} \end{cases}$$

The MPC is given by:

$$MPC = \frac{\Delta c_0}{\Delta y_0} = \frac{\tilde{c}_0(y_0 + \epsilon y_0) - c_0(y_0)}{\epsilon y_0} = \begin{cases} \frac{1}{1+\beta^{\frac{1}{\gamma}} R^{\frac{1-\gamma}{\gamma}}} & \text{if } \epsilon > 0 \\ \min \left\{ \frac{1}{1+\beta^{\frac{1}{\gamma}} R^{\frac{1-\gamma}{\gamma}}} \left(\frac{1+\epsilon}{\epsilon} \frac{1+\beta^{\frac{1}{\gamma}} R^{\frac{1-\gamma}{\gamma}}}{1+(\frac{\beta}{1-\lambda})^{\frac{1}{\gamma}} R^{\frac{1-\gamma}{\gamma}}} - \frac{1}{\epsilon} \right), 1 \right\} & \text{if } \epsilon < 0 \end{cases}$$

C.4 DA model with DA in t=0,1

Introducing DA in $t = 1$ yields the following optimization problem:

$$\begin{aligned} \max_{c_0, c_1, a_0^{plan}, a_1^{plan}} \quad & u(c_0) - \lambda d(a_0, a_0^{plan}) + \beta \left(u(c_1) - \lambda d(a_1, a_1^{plan}) \right) \\ \text{s.t.} \quad & c_0 + a_0 = y_0; \quad c_1 = R a_0; \end{aligned}$$

The savings plan is formed at the beginning of each period. We already know a_0^{plan} from the problem without DA in $t = 1$. The formation of a_1^{plan} is trivial as it is always optimal to consume everything in the final period, i.e. $a_1^{plan} = 0$. DA in $t = 1$ introduces two new Euler conditions, as marginal utility tomorrow now also depends on the difference between planned savings and actual savings tomorrow.

$$\begin{aligned} u'(c_0) &= \beta R u'(c_1) & \text{if } a_0 \geq a_0^{plan} \text{ and } a_1 \geq a_1^{plan} \\ (1-\lambda)u'(c_0) &= \beta R u'(c_1) & \text{if } a_0 < a_0^{plan} \text{ and } a_1 \geq a_1^{plan} \\ u'(c_0) &= \beta R(1-\lambda)u'(c_1) & \text{if } a_0 \geq a_0^{plan} \text{ and } a_1 < a_1^{plan} \\ (1-\lambda)u'(c_0) &= \beta R(1-\lambda)u'(c_1) & \text{if } a_0 < a_0^{plan} \text{ and } a_1 < a_1^{plan} \end{aligned}$$

Because it is always optimal to consume all savings in $t = 1$, and the optimal savings plan in $t = 1$ is always $a_1^{plan} = 0$, conditions three and four are irrelevant. Note that $a_1 < a_1^{plan} = 0$ would imply negative assets at death. As such, it is equivalent to the problem without DA in $t = 1$.

D Additional model results

Table D1: Model moments versus data moments

	Data	Model
Average wealth-to-income ratio	4.28	4.28
Average MPC out of losses	0.73	0.73
Ratio of households with savings plan/dissaving-aversion ratio between bottom and top quintile of wealth distribution	1.29	1.29

E Other Theories

This sections discusses briefly (and non-exhaustively) other common models of consumption behaviour in the context of MPC asymmetries.

E.1 Standard extensions

Most standard extensions generate higher MPCs (and MPC asymmetries) by generating a higher share of liquidity-constrained households, i.e. by shifting households along the wealth distribution. [Kaplan and Violante \(2022\)](#) provides an excellent overview of this literature. These extensions do not address, however, why unconstrained households responds asymmetrically to changes in income. In what follows, I discuss each extension in more detail.

Risk aversion: The effect of changes in risk aversion on MPC asymmetries is theoretically ambiguous and quantitatively small. Higher risk aversion concavifies the consumption function, but at the same time shifts households away from the borrowing constraint due to a stronger precautionary savings motive. These forces have off-setting effects on the MPC and MPC asymmetry and are therefore relatively small, as discussed in [Kaplan and Violante \(2022\)](#).

Discount-factor heterogeneity: Heterogeneity in the discount factor primarily affects the distribution of wealth, but has little bearing for MPC asymmetries. It shifts a larger share of (impatient) households closer to the borrowing constraint and generates a set of (patient) households with large wealth holdings. As such, it suffers from the missing-middle problem, i.e median wealth that's substantially below average wealth in excess of what the data suggest (as discussed in [Kaplan and Violante \(2022\)](#)). Apart from these shifts along the wealth distribution, discount factor heterogeneity only affects the level, but not the asymmetry of MPCs.

Return rate heterogeneity: Heterogeneity in returns generates similar results as discount-factor heterogeneity ([Kaplan and Violante, 2022](#)). Because a few high-return households hold

most of the wealth in the economy, this pushes down the discount factor that is necessary to match average wealth in the economy. A lower discount factor increases the MPC, but again neither introduces asymmetries for constrained nor unconstrained households.

Two-asset model: A model featuring both a liquid and illiquid asset in the spirit of [Kaplan and Violante \(2014\)](#) generates a larger share of liquidity-constrained households by introducing wealthy hand-to-mouth households. This roughly triples the share of liquidity-constrained households to one-third of the population ([Kaplan et al., 2014](#)). It does not provide an explanation for why the remaining two-thirds of households which are unconstrained would respond asymmetrically to changes in income. One could hypothesize that for the majority of households liquid resources are not sufficient to fully absorb the income loss, but this seems to be at odds with the data.

Consumption adjustment costs: [Fuster et al. \(2021\)](#) introduces a fixed utility cost to adjusting consumption and shows that it generates a larger sign asymmetry than a model without adjustment costs. However, this asymmetry is quantitatively small. The symmetric adjustment cost primarily addresses the extensive margin of consumption adjustment and increases the share of households that do not adjust consumption in response to a change in income. This affects both the response to positive and negative income shocks. A large shock of 10% of annual income is likely to be large enough to induce most households to pay the fixed cost and adjust consumption as the benefits of consumption smoothing outweigh the cost of consumption adjustment.

Asymmetric portfolio adjustment costs: Suppose there is a cost to adjusting your stock of assets that is asymmetric. Increasing your asset stock is costless, while decreasing your asset stock is costly. Irrespective of this cost being fixed or proportional to the adjustment, such a framework does not necessarily generate asymmetric MPCs. Even though liquidating assets in order to buffer income shocks is costly in this framework, the fraction of households *having to* liquidate assets will be too small to generate quantitatively meaningful MPC asymmetries. A large share of households are active savers, i.e. they save a fraction of income to accumulate wealth. Instead of liquidating wealth, they can simply reduce the fraction of income saved without having to pay the portfolio adjustment cost. For these households, the adjustment cost will have little bearing on the MPC asymmetry.

E.2 Behavioural explanations

Hyperbolic discounting: Hyperbolic discounting, as for example studied in [Laibson et al. \(2021\)](#), increases MPCs, but does not generate meaningful asymmetries as it amplifies both the response to gains and losses. In the simple two-period framework from Section 4, the MPC with

hyperbolic discounting would be given by:

$$MPC = \frac{1}{1 + \delta\beta}$$

where δ denotes the hyperbolic discount factor. Compared to a standard model, present bias increases MPCs by discounting future consumption at a higher rate, but it equally does so for gains and losses.

Rational inattention: [Reis \(2006\)](#) introduces a theory of inattentive households in which households only sporadically update their consumption plans due to a cost to processing information. Inattention introduces stickiness to consumption plans and lowers MPCs out of unexpected shocks, for both gains and losses. Alternatively, one can introduce households that plan savings instead of consumption. This increases MPCs out of both gains and losses as now most of the income change is absorbed through consumption instead of savings. Neither specification, sticky consumption plans nor savings plans generates meaningful asymmetries, however.

Temptation preferences: Temptation preferences introduce a demand for commitment and are similar in spirit to hyperbolic discounting with sophisticated agents. [Attanasio et al. \(2020\)](#) use such preferences to generate demand for illiquid assets. By locking away their wealth in housing, which is associated with a fixed cost, households can reduce the utility cost of temptation. As such, temptation preferences present one way to generate a large share of wealthy hand-to-mouth households without assuming excessively large returns on illiquid assets. Given that temptation preferences introduce an element of present bias, they tend to increase MPCs for both gains and losses, similar to hyperbolic discounting. Moreover, a larger share of wealthy hand-to-mouth households implies a larger share of households with asymmetric MPCs. But similarly to the two-asset model in [Kaplan and Violante \(2014\)](#), it does not explain why unconstrained households have asymmetric MPCs.

Expectations-based reference-dependence and loss aversion: [Pagel \(2017\)](#) studies a life-cycle model with expectations-based reference-dependent preferences, building on previous work by [Kőszegi and Rabin \(2009\)](#). Within this framework, household's period utility not only consists of the standard utility from consumption, but additionally of gain-loss utility, i.e. the deviation of consumption relative to a reference point. This reference point is given by the household's previous expectations about both present and future consumption. The preference structure generates excess smoothness and sensitivity in consumption. However, because unexpected losses in present consumption are more painful than expected losses in the future, households delay unexpected losses in consumption until expectations have adjusted in the future. This lowers the MPC out of losses and therefore does not address the MPC asymmetry found in this paper.